When a fresh screen is not enough: safe retry timing for computer-use agents

Date: 2026-10-15 (America/Toronto). Research type: formal analysis, exhaustive synthetic event study, and controlled real-agent browser pilot. Codex generated this paper, code, and analysis. No human or external peer review occurred. The added pilot tests real agent decisions in a local browser task.

Abstract

Computer-use systems can repeat an action when its result is not visible. A current empty read does not show that an earlier action can no longer take effect. We compare retry rules in two separate studies. The finite synthetic study covers 336 constructed cases and 13,440 runs. In its main 256-case append grid, first-empty retry creates 60 duplicates with checks every tick, versus 36 with checks every four ticks. Truthful terminal status and a conservative timing bound both produce zero duplicates and 192 clean confirmations. The terminal rule reduces matched synthetic mean time from 9.500 to 7.333 ticks. A later pilot uses two real Codex agent contexts in Chrome for 18 controlled local trials. Absence retry creates four duplicates in six trials. Terminal and a 24-second timeout each create zero duplicates and five clean on-time claims in six trials. One terminal claim is late. The four clean timing pairs include two faster and two slower terminal trials; they do not establish a general browser speed gain. All four browser duplicate claims occur after two records already exist. No browser effect occurs after stopping. The contribution is a repeatable comparison that separates read freshness, execution closure, claim timing, and final effects. It applies established systems principles. The browser task is constructed, and its counts are not field error rates.

1. Research question and contribution

An agent clicks a button. The next observation shows no new item. Should the agent click again? This question applies to any action that creates an additional record when repeated. A **non-idempotent action** can change the intended result when it is repeated. An **idempotent action** has the same intended effect when repeated [5].

We ask how the completion delay of an earlier action and the delay of an observation affect retry safety, confirmed completion, and observation cost. **Execution closure** means that an earlier attempt can no longer create a future effect. A current read and execution closure are different properties.

The contribution of this paper is its specific test design and full comparison. We combine pending writes, delayed negative reads, observation schedules, and post-stop effect checks. We include a conservative no-retry control, an adequate fixed-bound baseline, and a matched server-deduplication control. We publish every event trace, including writes that finish after the controller stops. This contribution is a small synthetic benchmark, interface audit, and controlled browser pilot. It is not a new exactly-once protocol, learned agent, or general theorem. The limited source search cannot prove that this is the first such comparison.

Reliable completion matters because an extra record can be harmful even when the original task eventually succeeds. Faster tools are useful only if the controller interprets their observations correctly. We therefore keep duplicate prevention, useful completion, delay, and call count separate.

2. Related work

CUA-Speedrun studies speed and cost across computer-use configurations. Its fast input-output analysis shows that a screenshot can arrive before an application updates its interface. The tested agent can then take extra steps. Our study isolates retry effects with controlled policies. It does not reproduce that model experiment [1].

RStackBench separates visible success from persisted task state. Its controlled browser tests distinguish delayed persistence from a missing required action. It tests individual fault layers and identifies combined faults as future work. Our model combines a pending non-idempotent write with a negative read and scores all later effects. We do not reuse its code, tasks, or data [2].

CapLease studies repeated agent proposals after uncertain execution. It uses durable authorization state and states the need for a suitable external idempotency contract. We do not propose a replacement for that system. Our comparison concerns the evidence that permits one retry, under explicitly different server contracts [3].

RIFL already treats lost results and duplicate execution as systems problems. Its browser example discusses uncertain outcomes and recognition of prior effects. HTTP Semantics also limits automatic retries of non-idempotent requests. These sources establish that absence of a visible result is not a general retry guarantee [4,5].

The original source record contains 15 exact queries, read sections, versions, access limits, and additional related work. The browser extension adds two queries and a new check of the closest sources. RStackBench already reports a small model planning pilot; real-agent execution alone is not a new research idea [7]. See source_search.md, browser_source_search.md, and references.bib. The study does not claim that waiting, verification, durable status, or duplicate suppression is new.

3. Method

3.1 Saved plan and finite environment

The plan and input configuration were saved before implementation or numerical analysis. A design review then changed the completion guard, still before analysis: an observed count of two stops with a detected duplicate and no completion claim. An observed count of one is the only completion claim. The original hash, amendment time, and final hashes are in input_manifest.json. No cases or hypotheses changed.

Time is an integer **tick**, a synthetic time unit. It is not a measured second. The initial attempt starts at tick zero. Each attempt resolves after its assigned delay, either by committing one durable append or by dropping permanently with no effect. At most one retry is allowed. Both attempts refer to one desired item. A dropped attempt cannot later commit.

At time t, a query returns the exact count and attempt status at historical time max(0, t-L), where L is the case's fixed read lag. Status is pending, committed, or permanently dropped. Due resolutions occur before a query at the same tick. The retry is sent after the query. Each attempt has a positive delay. No controller sees the future delay or outcome before the observation reveals it.

The controller stops on a positive count: one gives a completion claim; two gives a detected duplicate. With no positive result it stops at tick 24. A deadline stop does not cancel an admitted action. Every admitted action continues until resolution, including actions sent at the deadline. The final audit occurs after the last admitted resolution or controller stop, whichever is later.

The model has no other writers, rollback, incorrect status, identity error, read error, or crash. Query calls have no added duration. Delay and call count are therefore separate measures, not a full service-latency model.

3.2 Five rules and two server contracts

Each rule makes the same initial submission and follows the same chosen query schedule. A retry is possible only after an empty observation and only if none has been sent before.

Rule	Condition that permits the one retry
No retry	Never retry
Absence retry	The first empty observation
Timeout 8	An empty observation at tick 8 or later
Bound 12	An empty observation at tick 12 or later
Terminal status	The first attempt is observed as permanently dropped

Timeout 8 uses the main maximum execution delay but ignores read lag. Bound 12 adds the maximum execution delay of 8 to the maximum lag of 4. It is an adequate timing baseline under the main contract. The terminal rule uses information from the same response that the other rules ignore. This extra information is an application capability, not an ability inferred from a screenshot.

Under **append semantics**, each successful attempt creates another item. Under **ideal deduplication**, all successful attempts for the same logical item produce at most one item. A dropped attempt does not prevent a later attempt from succeeding. This matched positive control assumes correct stable identity and server enforcement. The client cannot create this guarantee by assigning a local identifier alone.

3.3 Cases, schedules, and measures

The main grid uses both attempt delays in {1,2,4,8}, both outcomes in {commit,drop}, and lag in {0,1,2,4}. This gives 256 cases. The four query schedules are every tick, every 2 ticks, every 4 ticks, and backoff at {1,2,4,8,12,16,20,24}. Five rules and two server modes give 10,240 main runs.

The late-execution stress suite sets both delays to 12 and retains the four main lags: 16 cases. The stale-read suite uses the main delays but lag 8: 64 cases. Both use all rules, schedules, and modes. There are 3,200 stress runs. These suites break one timing bound and remain separate from the main grid.

A **duplicate** is a final effect count greater than one. A **clean confirmation** is a completion claim by tick 24 with exactly one final effect. An **unsafe claim** observes one item, claims completion, and later ends with a duplicate. A detected duplicate and a deadline stop make no completion claim. We also save zero effects and one effect without confirmation.

Mean clean time is the mean confirmation tick among clean cases. Comparisons between terminal status and Bound 12 use their jointly clean cases. The **capped clean-time score** uses the confirmation tick for a clean result and 24 otherwise. This defined score is not wall time. Query and submission counts are reported separately. Full rows and grouped summaries preserve all measures.

All cases are constructed. Equal weighting is a description of this grid, not an estimate of application frequencies. Some delay values repeat when a branch is not taken. These are designed factor combinations, not independent experimental observations. There is no training set, random sampling, fitted parameter, seed, significance test, or population confidence interval.

4. Formal boundary

A simple trace shows why freshness is insufficient. At tick 1, a current read sees zero items while the original attempt is pending. A retry commits at tick 2. The controller sees one item and stops. The original attempt then commits at tick 8. The final count is two. Scoring only the controller's stop state would miss the failure.

Consider two worlds with the same observation history before a retry decision. In one, the original attempt will commit later. In the other, it will drop. Suppose the retry succeeds and there is no later distinguishing observation before the decision. If a policy retries with probability r, its duplicate probability in the commit world is r. Its recovery probability in the drop world is also r. Thus a duplicate limit epsilon on this local decision also limits recovery to epsilon. This is a tight, conditional

two-world result. It does not apply after new distinguishing evidence arrives and is not a claimed new distributed-systems theorem.

For terminal status, a retry is sent only after a truthful permanent drop. The original attempt therefore cannot add an effect later. At most one remaining attempt exists, so the final count is at most one. Delayed visibility can delay this decision but cannot reverse a permanent drop.

For Bound 12 in the main model, a query at t >= 12 has an as-of time at least 8. The first attempt has resolved by tick 8. If the returned count is zero, that attempt did not commit and cannot do so later. One retry is then safe. This argument fails when the declared execution or lag bound is false. Full assumptions, proofs, and the local ambiguity check are in derivations.md.

5. Results

5.1 Current negative reads still permit unsafe retries

The following results use append semantics and queries every tick. Each rule has 256 cases. Time is in ticks. Mean clean time must not be compared across rows with different successful case sets as if those sets were matched.

Rule	Duplicates	Clean confirmations	Unsafe claims	Mean queries	Mean clean time
No retry	0	128	0	14.750	5.500
Absence retry	60	132	56	10.109	5.848
Timeout 8	12	180	12	12.125	8.022
Bound 12	0	192	0	13.125	9.500
Terminal status	0	192	0	11.500	7.333

Absence retry creates 60 duplicates (23.44% of this grid). Of these, 56 occur after a one-item completion claim; four are already visible when the controller stops. These are different failure records. Both count as duplicate effects.

In the 64 zero-lag cases, absence retry still produces 12 duplicates (18.75%). This supports H1: current state can be correct while the earlier operation is unfinished. The trace above is one such case. All five policies for that case are in trace_examples.json.

No retry avoids duplicates but confirms only 128 cases. Terminal status and Bound 12 both confirm 192, with zero duplicates. The remaining 64 cases drop both possible attempts. Thus the safe rules recover every case in this main grid where at least one of the two possible attempts can commit. This is a property of the designed cases and deadline, not a real-world success rate.

5.2 Terminal status improves the safe timing baseline

Terminal status and Bound 12 have the same 192 clean cases for every schedule. With queries every tick, mean time falls from 9.500 to 7.333 ticks, a reduction of 2.167 ticks or 22.81%. Mean queries across all 256 cases fall from 13.125 to 11.500, a 12.38% reduction. Both use 1.500 submissions per case. Their capped clean-time scores are 13.125 and 11.500 ticks. H2 holds in all four tested schedules.

The improvement comes from earlier evidence that the first attempt permanently dropped. It does not come from hidden knowledge of the outcome. A service without this status signal cannot assume the same benefit.

Query schedule	Absence duplicates	Timeout 8 duplicates	Bound 12 clean time	Terminal clean time
Every tick	60	12	9.500	7.333
Every 2 ticks	52	12	9.875	7.833
Every 4 ticks	36	12	11.000	9.333
Backoff	60	12	10.708	8.885

5.3 Faster checks can increase duplicate effects

For absence retry, changing from every four ticks to every tick adds 24 duplicate cases: 36 to 60, or 9.375 percentage points in this grid. Faster reads also change which duplicates are detected before stopping. The paired counts are 24 cases worse with every-tick queries, zero worse with every-four-tick queries, 36 with duplicates in both, and 196 with duplicates in neither. This supports the planned directional comparison H3. It does not mean that slow tools are generally safer or better.

The safe rules show a different cost tradeoff. Terminal status uses 11.500 queries per case with every-tick observation and 3.250 with every-four-tick observation. Its mean clean time rises from 7.333 to 9.333 ticks. Backoff uses 5.094 queries and 8.885 clean ticks. The study does not select an optimal schedule or attach a monetary value to a query.

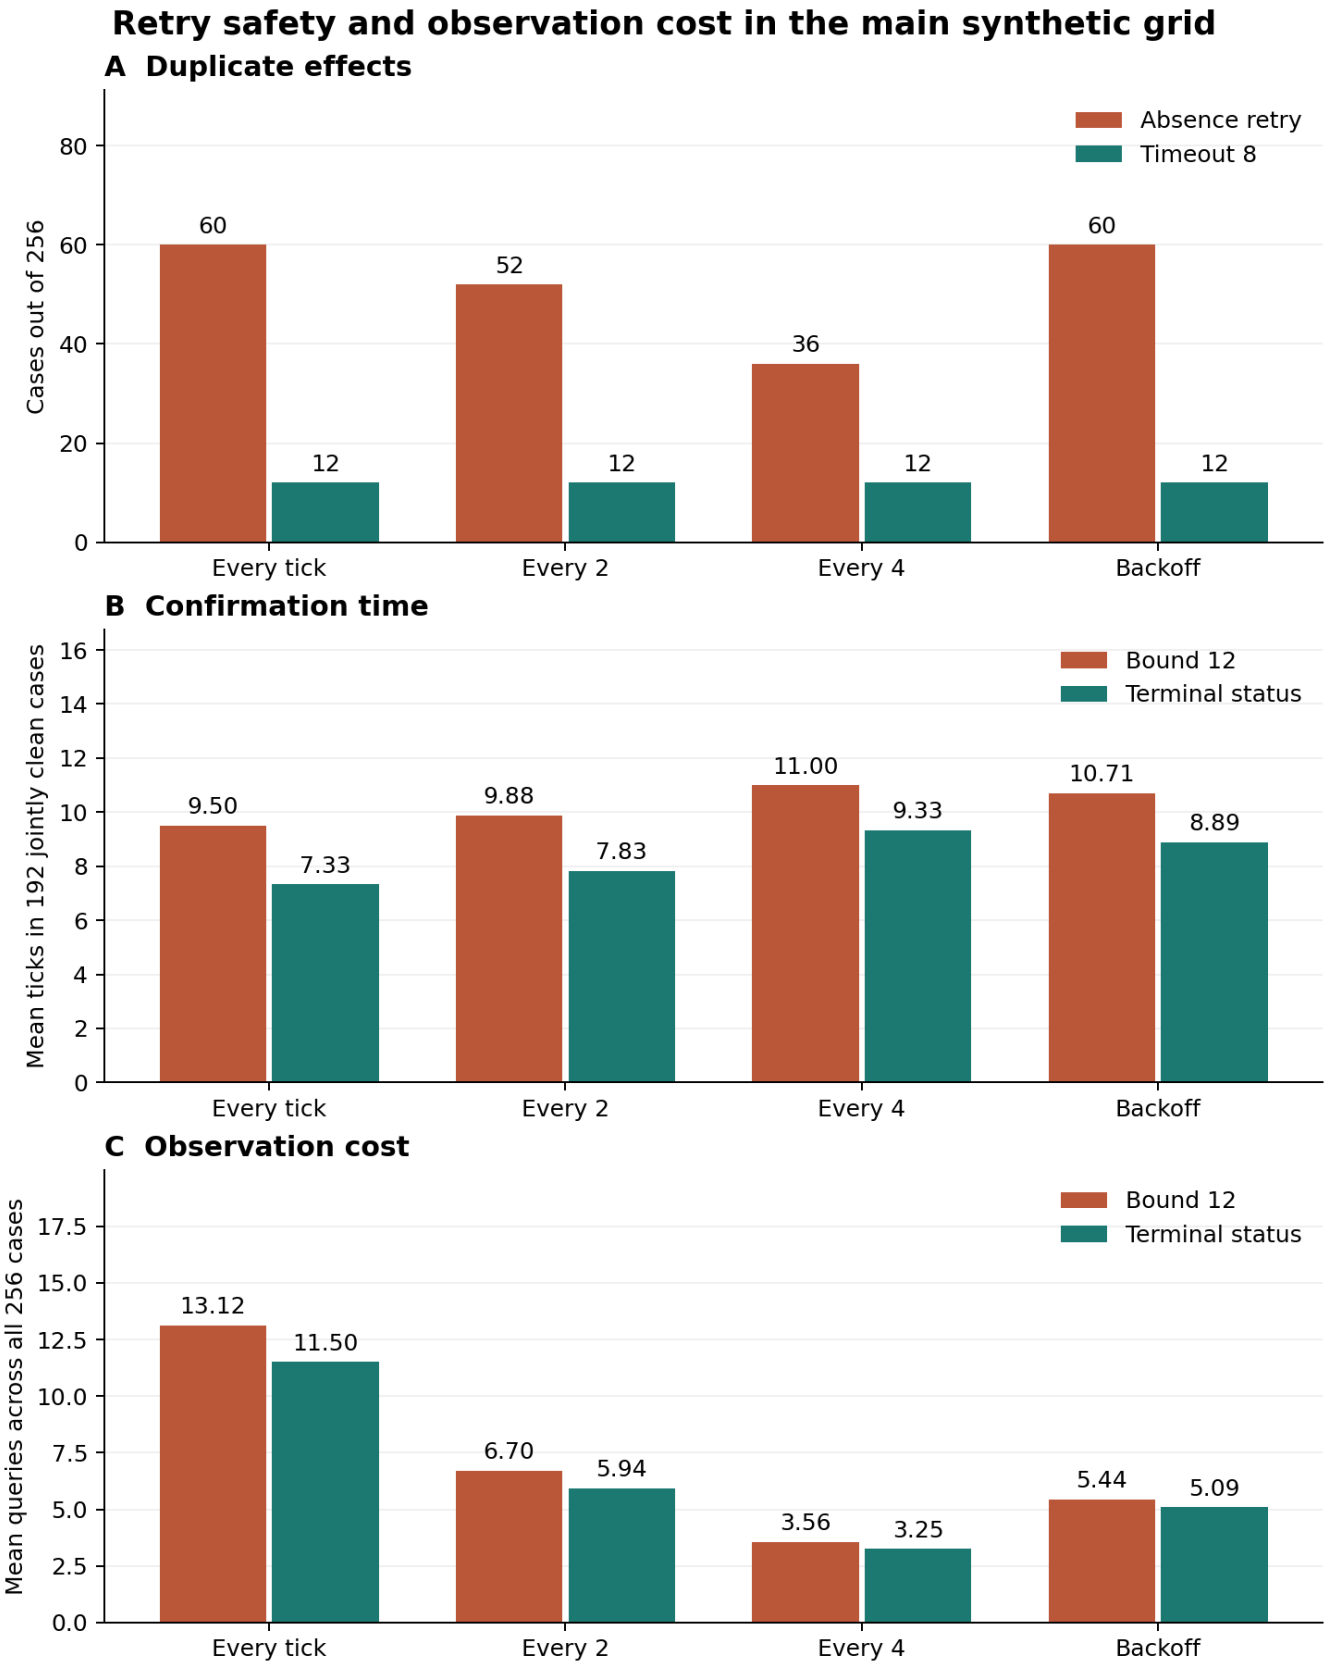


Figure 1. Main synthetic append grid. Panel A uses all 256 cases. Panels B and C compare terminal status with Bound 12. Their clean case sets match. Panel B uses those 192 cases; Panel C uses all 256. Neither safe rule has a duplicate in this grid. Backoff is the fixed schedule in Section 3.3.

5.4 Bound violations and server deduplication

These stress results use every-tick queries. They must not be pooled with the main cases.

Bound violation	Rule	Cases	Duplicates	Clean confirmations	Deadline stops
Delay 12	No retry	16	0	8	8
Delay 12	Bound 12	16	3	6	7
Delay 12	Terminal status	16	0	9	7
Read lag 8	No retry	64	0	32	32
Read lag 8	Bound 12	64	4	40	20
Read lag 8	Terminal status	64	0	41	23

The 12-tick rule creates three duplicates in the late-execution suite and four in the stale-read suite. The terminal rule creates none. However, it confirms only 9 of 16 and 41 of 64 cases before the deadline. Safety does not imply timely completion. The full schedule results are retained even when they are worse.

Under ideal server deduplication, no policy produces a duplicate in any of the 6,720 matched runs. This confirms the intended positive control. It does not demonstrate an implemented or universal exactly-once service. H4 holds under the stated truthful-status and deduplication contracts.

6. Real agent and browser pilot

6.1 Saved design and executed task

We added a controlled browser pilot after the synthetic study. The new plan and configuration were saved before implementation and scored execution. The implementation was hashed before the warm-up and scored trials. The original synthetic results remain unchanged. See browser_plan.md, browser_plan_manifest.json, and browser_execution_manifest.json.

Two separate Codex agent contexts used Chrome through the supported computer-use interface. Each context received nine trials in a fixed order. Runner B used the reverse policy-condition order of runner A. The agents read the rendered accessibility or DOM state and used visible buttons. The runners were instructed to select each next action in a separate model decision and not use a scripted policy loop. The saved logs record actions and states, but do not contain model-turn boundaries or tool code. Thus, those logs cannot alone prove every access or decision constraint. This is a real agent interaction test with an artificial local task. It is not a pixel-only benchmark or a public-web evaluation.

The task was to create exactly one stored item. A local Python HTTP server accepted submissions. Independent timers then committed or permanently dropped each attempt. Each commit inserted a row into SQLite. A requested state check showed the item count and first-attempt status at a configured historical time. That snapshot did not update by itself. The agent could submit, request state, claim completion, or stop. The app enforced a maximum of two submissions. It did not enforce the policy-specific retry or claim rules. The state fields were structured and truthful by construction.

Condition	Initial attempt	Retry if sent	Read lag
Pending	Commit after 16 seconds	Commit after 2 seconds	0 seconds
Failure	Permanently drop after 4 seconds	Commit after 2 seconds	0 seconds
Stale	Commit after 4 seconds	Commit after 2 seconds	8 seconds

These are nominal timer delays. The evidence records actual resolution times. Three rules each had two runs per condition, for 18 assigned trials. Absence retried after the first requested empty state. Bound24 required a newly requested empty state at elapsed time 24 seconds or later. Terminal required a requested empty state with permanent failure status. All rules allowed one retry at most.

The 24-second threshold adds the largest nominal action delay and largest read lag. It is conservative for these selected combinations. It is not the fastest possible safe timeout.

A requested count of one permitted a completion claim before the 60-second deadline. A count of two required a duplicate stop. A state check at or after 60 seconds required an unresolved stop. A late claim was retained but could not count as a clean on-time confirmation. Pending actions continued after every stop. Final stored rows, not the stop screen, determined duplicate outcomes.

The runner prompts hid condition labels, delay values, outcomes, source code, and audit data. The runners were told not to inspect those data or the other runner. They shared one machine and browser service, so this was instruction-based separation, not a secure isolation boundary. The lead did one excluded warm-up trial. The fixed budget allowed new scored trials for 12 minutes after the first scored submission. There was no replacement of an inconvenient trial.

6.2 Measurements and limits of comparison

The primary measures were final duplicates, clean on-time claims, unsafe claims, and adherence to the supplied rule. An unsafe claim is any completion claim whose final item count is not one. The server used a monotonic clock for all measured elapsed seconds. Browser and model overhead contributed to these times. The study did not measure their parts separately.

We retained each browser action log, requested snapshot, server event, durable record, and stopped-trial screenshot. The final audit waited for all accepted attempts to resolve. A negative observation was called stale when a committed record already existed at the observation request time but was absent from its historical snapshot. An unresolved-negative exposure was an empty observation made before the initial attempt resolved. A pending-status exposure could also come from a delayed historical snapshot. These are test exposures, not model mistakes by themselves.

The saved hypotheses asked whether absence could create duplicates, whether terminal could avoid them and recover earlier than Bound24, and whether the final audit could find a duplicate after a completion claim. All assigned trials remain in the results. Timing comparisons pair the same runner and condition where both policies have clean on-time claims. Two repetitions per cell do not support a population error rate, a safety certificate, or a statistical significance claim. We do not combine these seconds or counts with the synthetic grid.

6.3 Observed results

All 18 assigned trials started and stopped. The frozen evidence contains 380 server events, 23 stored rows including the excluded warm-up, and 62 requested snapshots including the warm-up. The scored trials used 28 accepted submissions and 61 snapshots. The browser reported Chrome 155.0.0.0 on macOS. The server used CPython 3.14.7. The exact immutable agent model build was not available. See browser_environment.json.

Rule	Trials	Duplicates	Clean on-time claims	Unsafe claims	Deadline misses	Snapshots
Absence	6	4	1	4	1	13
Bound24	6	0	5	0	1	30
Terminal	6	0	5	0	1	18

Deadline misses include unresolved stops and late completion claims. Absence had one unresolved stop. Bound24 had one unresolved stop. Terminal had one late claim. The four duplicates came from both pending and both stale-condition absence trials. Three had an empty observation before a later original commit. All three created duplicates, which meets the narrow H-B1 criterion. The fourth, A09, saw an empty historical state after the original record already existed. It also created a duplicate. Thus both unresolved execution and delayed visibility were active in the observed failures. All four false completion claims were allowed by the deliberately weak absence rule. They are not evidence that an unconstrained agent would naturally choose that rule.

A01 gives a measured example with no configured read lag. The first empty check occurred at 4.091 seconds. The agent submitted a retry at 7.087 seconds. That retry committed at 9.097 seconds. A check at 12.560 seconds showed one item while the original attempt remained pending. The original attempt committed at 16.008 seconds. The completion claim arrived at 16.244 seconds and referred to the earlier one-item snapshot. The final database held two items. The other three duplicate trials also had two stored items before the claim reached the server. There were no post-stop effects and no duplicate created after a claim. H-B3 therefore had zero observed instances. The pilot does not reproduce the synthetic study's later-after-stop duplicate trace.

Bound24 and Terminal each had five clean on-time claims and no duplicates. Terminal used 18 snapshots, versus 30 for Bound24, across six assigned trials each. These are observed call totals. Different delays, interrupted contexts, and outcome sets prevent a causal efficiency claim. The server-visible rule checks passed in 17 trials. A04 made a late terminal-policy claim at 69.230 seconds after reading a one-item snapshot requested at 20.076 seconds. Its final count was one, but it failed the on-time criterion. This is a deadline failure, not a duplicate failure.

The agent logs report context recovery or compaction delays in A02, A04, B05, and B07. A02 stopped unresolved at 93.147 seconds. B05 stopped unresolved at 68.737 seconds. Both had one final item. B07's first check arrived at 49.340 seconds; its clean claim arrived at 52.488 seconds. Repeated Statsig telemetry HTTP 403 notices appeared in both runner logs. The saved task requests reached the local server. The study cannot isolate how much time came from telemetry, model inference, tool work, context recovery, or shared host load. No trial was removed or repeated.

6.4 Paired timing results

The table preserves all six planned pairs. The difference is Terminal minus Bound24, so a negative value means the terminal trial was faster. A difference is given only when both trials were clean and on time. All values are measured seconds.

Runner and condition	Bound24 claim	Terminal claim	Paired difference	Pair status
A, failure	None	17.789	Not compared	Bound24 unresolved
A, pending	25.488	69.230	Not compared	Terminal late
A, stale	21.268	16.529	-4.739	Both clean
B, failure	37.583	22.092	-15.491	Both clean
B, pending	21.599	23.272	+1.673	Both clean
B, stale	15.364	52.488	+37.124	Both clean

The four matched differences range from -15.491 to +37.124 seconds. Their median is -1.533 seconds and mean is +4.642 seconds. Terminal is faster in two pairs and slower in two. Only the B failure pair is clean under both rules. In that pair, Terminal confirms 15.491 seconds earlier than Bound24. The A failure pair is not a valid speed comparison because Bound24 has no completion claim. H-B2 supports zero duplicates under the tested truthful-status contract, but these browser data do not establish a general speed gain. In particular, the synthetic 22.81% reduction must not be presented as a browser measurement.

No trial breached the selected 24-second first-resolution-plus-lag threshold. The largest timer overrun above its nominal scheduled delay was 0.012 seconds. Nominal timer values were not exact real-time guarantees. The saved raw times retain full precision. The 60-second claim deadline and tool delay had larger effects on clean completion than these small scheduling overruns.

The full trial table, all six pairs, and evidence checks are in browser_results.csv, browser_pairs.csv, and browser_summary.json. The pilot supports the specific duplicate-risk mechanism and shows a

separate deadline problem. It does not establish field error rates, natural agent behavior, or universal safe retry timing.

7. Computer-use interface implications

The results support a concrete interface requirement: distinguish an unresolved attempt from a terminal failure with no future effect. A timeout, unchanged screen, or current empty record list is not that failure signal. Where the application supports it, bind status to the specific attempt and keep that identity across recovery.

A controller should retain pending actions after its own wait ends. It should use a documented terminal signal, a valid execution-and-visibility bound, or a server-enforced idempotency contract before a repeat action. If none is available, the outcome remains unresolved. Waiting or asking for review has an availability cost; it does not prove failure.

An evaluator should also let admitted actions finish after the agent stops. In this grid, that step exposes 56 unsafe early claims for every-tick absence retry. A test that checks only the visible count at stop can award success before a duplicate appears. The browser late claim also shows why a deadline should be checked at claim receipt. A snapshot timestamp can be valid while the later claim is too late. These recommendations apply known systems reasoning to computer-use interfaces. The controlled browser pilot tests the same distinction with real requests and writes. Public-web tasks and broader model comparisons remain future work.

8. Limitations

The main grid runs a finite event model. Its tick values are chosen; its call costs and delays are not measured on a service. The browser extension uses real model decisions, requests, timers, and stored writes, but its task and fault conditions are constructed. The grid weights are artificial. Real applications can have unbounded delays, rollback, partial writes, status loss, incorrect status, changing caches, and multiple writers.

The terminal rule assumes a truthful, durable, attempt-specific drop record. Some sites expose no such record. The model gives all rules the same response but only this rule uses that extra field. Its advantage is therefore conditional on an interface capability. A richer failure model can invalidate the guarantee.

Only one retry and one logical item are modeled. There is no crash recovery, concurrent controller, permission change, authentication failure, or imperfect action identity. Ideal deduplication is a control, not a tested implementation. The deadline is fixed, query duration is zero, and the schedules are not learned or optimized. Neither the finite cases nor the proof justify a safety certification or field error estimate.

The same cases are disclosed to the policy designers. The browser runners did not receive the condition configuration, but each received a prescribed policy. This does not measure natural retry behavior. There are only two contexts, one machine, one browser service, and one simple local task. Repeated trials share context within each runner. Tool and model delay can remove an intended negative-state exposure before a check occurs. The exact model build and service timing are not available. No held-out public-web task supports general transfer. Further work should test more agents, richer applications, and permitted real workflows.

9. Ethics, rights, and AI use

The study can help prevent duplicate records and premature completion claims. It contains no harmful prompts, exploit steps, private records, account data, or real transactions. None of the experiment sends messages or changes an external account. GitHub publication is a separate authorized action, not experimental evidence.

Codex generated the plan, implementations, calculations, figure, and paper. Two separate Codex contexts performed the browser trials. Other Codex agents checked sources, design, calculations, and reproduction. AI review is not human or external peer review. No human review, university affiliation, endorsement, or certification is claimed. The browser pilot used real model inference through the existing Codex service. No model training, separate paid API, or new paid compute service was used.

All input cases and local browser tasks are newly constructed. No source paper, benchmark task, dataset, code, or figure is redistributed. Sources retain their rights. No new license is assigned to this work. The paper uses original prose and cites the related methods.

10. Checks and reproduction

The main implementation advances the environment one tick at a time. The independent validator uses direct event-time thresholds and does not import the analyzer. It checked all 13,440 run rows, 167,520 trace events, 676 post-stop resolutions, 6,720 matched server pairs, all grouped result tables, and 48 paired comparisons. All checked integer and label fields match exactly. See method_review.md and validation_results.json.

The browser analyzer reads the frozen SQLite database and audit independently of the application code. It checks event, snapshot, stop, and stored-row identities before it computes outcomes. Its author also reviewed the design, so a fresh reviewer separately checks the final evidence. See browser_review.md and browser_reproduction_check.json. The action logs and screenshots support the observed interaction sequence; they do not reveal internal model reasoning.

Run the documented commands in README.md. The package includes the saved plan, cases, complete compressed traces, readable examples, summaries, code, dependency versions, logs, proofs, source record, and agent reviews. Relevant parts of the NeurIPS checklist were used to review claims, assumptions, reproduction, limits, rights, and AI use [6]. This is not a conference submission or approval claim. See final_review.md for the final reproduction outcome.

References

1. Aggarwal, P., Jang, L. K., Welleck, S., Fried, D., Salakhutdinov, R., and Koh, J. Y. (2026). *cua-speedrun: Standardized Benchmarking of the Speed of Computer-Use Agents*. arXiv:2609.40284v1, September 30. Preprint. Source.
2. Calanza, E. (2026). *RStackBench: A Five-Layer Fault-Injection Method for Measuring Browser-Agent Execution Reliability in Controlled Applications*. Independent author preprint, August. Source.
3. Xu, J., Fan, L., Wang, Z., Li, X., and Liu, H. (2026). *Beyond Single-Use Tokens: Durable Authorization State for Replay-Resistant LLM Agent Actions*. arXiv:2608.01710v1, August 3. Preprint. Source.
4. Lee, C., Park, S. J., Kejriwal, A., Matsushita, S., and Ousterhout, J. (2015). *Implementing Linearizability at Large Scale and Low Latency*. SOSP 2015. DOI: 10.1145/2815400.2815416. Author-hosted paper.
5. Fielding, R. T., Nottingham, M., and Reschke, J., editors (2022). *HTTP Semantics*. RFC 9110, Section 9.2.2. Internet Standard. Source.
6. NeurIPS. *Paper Checklist Guidelines*. Accessed October 15, 2026. Review aid.
7. Calanza, E. (2026). *Measuring the Gap Between Visible UI Success and Persisted State in Browser Agents*. Author report, August 14. Source.